


Load & Go for Db2 Warehouse on the IBM Cloud

—
Miran Badzak
Senior Product Manager

think 2018



Think 2018 / DOC ID / Month XX, 2018 / © 2018 IBM Corporation



Please note

IBM's statements regarding its plans, directions, and intent are subject to change or withdrawal without notice and at IBM's sole discretion.

Information regarding potential future products is intended to outline our general product direction and it should not be relied on in making a purchasing decision.

The information mentioned regarding potential future products is not a commitment, promise, or legal obligation to deliver any material, code or functionality. Information about potential future products may not be incorporated into any contract.

The development, release, and timing of any future features or functionality described for our products remains at our sole discretion.

Performance is based on measurements and projections using standard IBM benchmarks in a controlled environment. The actual throughput or performance that any user will experience will vary depending upon many factors, including considerations such as the amount of multiprogramming in the user's job stream, the I/O configuration, the storage configuration, and the workload processed. Therefore, no assurance can be given that an individual user will achieve results similar to those stated here.

Notices and disclaimers

© 2018 International Business Machines Corporation. No part of this document may be reproduced or transmitted in any form without written permission from IBM.

U.S. Government Users Restricted Rights — use, duplication or disclosure restricted by GSA ADP Schedule Contract with IBM.

Information in these presentations (including information relating to products that have not yet been announced by IBM) has been reviewed for accuracy as of the date of initial publication and could include unintentional technical or typographical errors. IBM shall have no responsibility to update this information. **This document is distributed “as is” without any warranty, either express or implied. In no event, shall IBM be liable for any damage arising from the use of this information, including but not limited to, loss of data, business interruption, loss of profit or loss of opportunity.** IBM products and services are warranted per the terms and conditions of the agreements under which they are provided.

IBM products are manufactured from new parts or new and used parts. In some cases, a product may not be new and may have been previously installed. Regardless, our warranty terms apply.”

Any statements regarding IBM's future direction, intent or product plans are subject to change or withdrawal without notice.

Performance data contained herein was generally obtained in a controlled, isolated environments. Customer examples are presented as illustrations of how those customers have used IBM products and the results they may have achieved. Actual performance, cost, savings or other results in other operating environments may vary.

References in this document to IBM products, programs, or services does not imply that IBM intends to make such products, programs or services available in all countries in which IBM operates or does business.

Workshops, sessions and associated materials may have been prepared by independent session speakers, and do not necessarily reflect the views of IBM. All materials and discussions are provided for informational purposes only, and are neither intended to, nor shall constitute legal or other guidance or advice to any individual participant or their specific situation.

It is the customer's responsibility to insure its own compliance with legal requirements and to obtain advice of competent legal counsel as to the identification and interpretation of any relevant laws and regulatory requirements that may affect the customer's business and any actions the customer may need to take to comply with such laws. IBM does not provide legal advice or represent or warrant that its services or products will ensure that the customer follows any law.

Notices and disclaimers continued

Information concerning non-IBM products was obtained from the suppliers of those products, their published announcements or other publicly available sources. IBM has not tested those products about this publication and cannot confirm the accuracy of performance, compatibility or any other claims related to non-IBM products. Questions on the capabilities of non-IBM products should be addressed to the suppliers of those products. IBM does not warrant the quality of any third-party products, or the ability of any such third-party products to interoperate with IBM's products. **IBM expressly disclaims all warranties, expressed or implied, including but not limited to, the implied warranties of merchantability and fitness for a purpose.** The provision of the information contained herein is not intended to, and does not, grant any right or license under any IBM patents, copyrights, trademarks or other intellectual property right.

IBM, the IBM logo, ibm.com and [names of other referenced IBM products and services used in the presentation] are trademarks of International Business Machines Corporation, registered in many jurisdictions worldwide. Other product and service names might be trademarks of IBM or other companies. A current list of IBM trademarks is available on the Web at "Copyright and trademark information" at: www.ibm.com/legal/copytrade.shtml.

Contents

- About Db2 Warehouse on the IBM Cloud
- Options for loading your data
 - Small loads via the built-in console
 - Pulling data from our IBM Cloud Object Storage with external tables
 - Pulling data from AWS S3
- Over-the-wire migration with Lift CLI
- Ship-disk solution for large data volumes
- Lift to Cloud with Integrated Analytics System



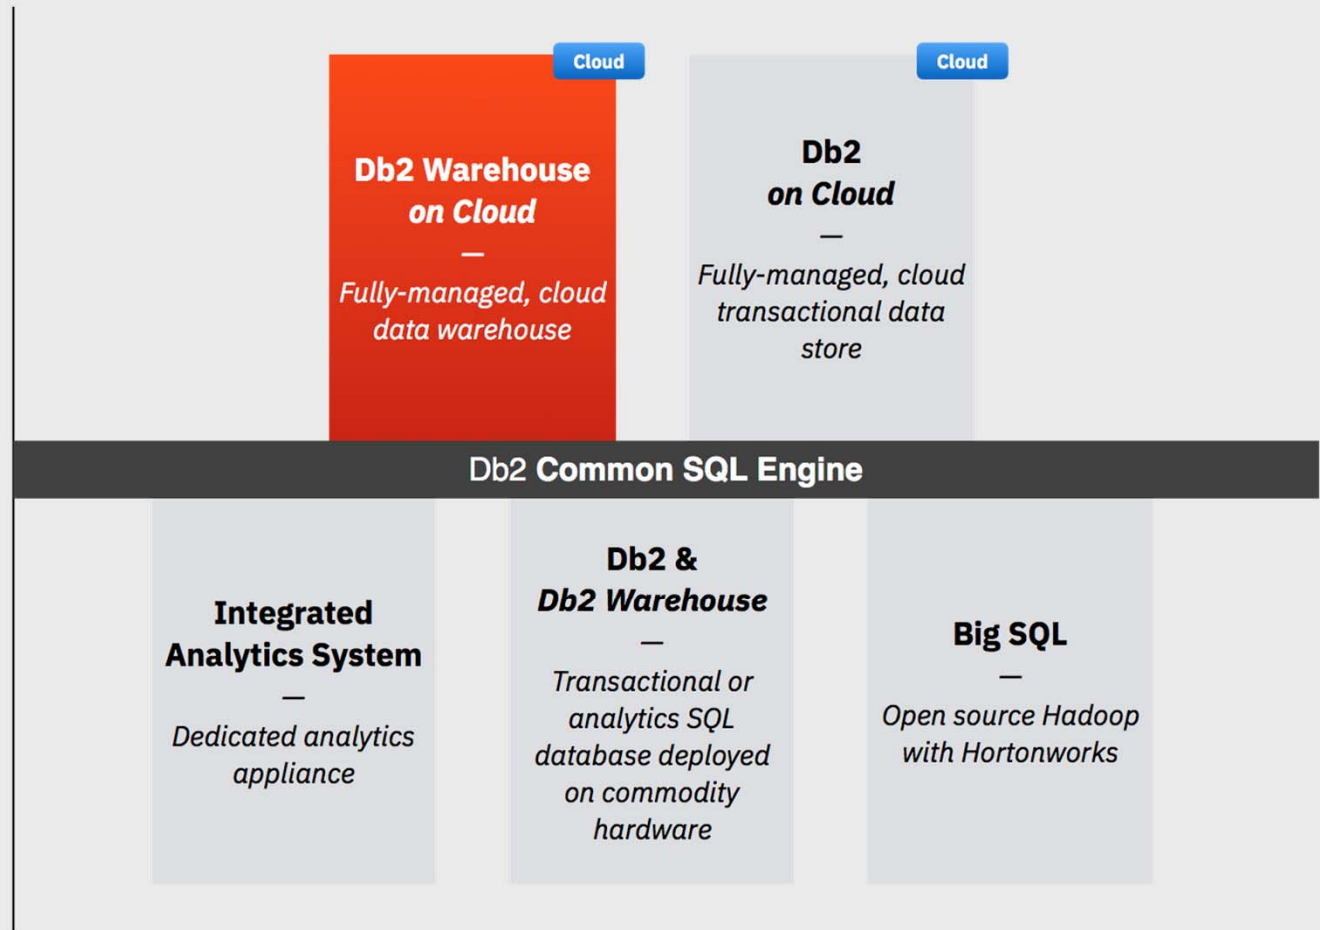
Db2

IBM Db2 Warehouse on Cloud is a fully-managed, high performance, petabyte-scale cloud data warehouse.

It delivers true elasticity with independent scaling of storage and compute, a highly-optimized columnar data store, actionable compression and in-memory processing, all working together to supercharge your analytics workloads.

Our family of
**Hybrid Data
Management**
solutions
built on the **Db2**
SQL engine

*Write your SQL once
deploy against any form factor
run anywhere*



Db2 Warehouse on Cloud

Fully-managed

by our world-class
DevOps team,
24x7x365

Blazing fast analytics

with IBM BLU
acceleration

Multiple architectures

to fit your performance
requirements and your
budget

And highly-compatible with your existing Netezza workloads
98%+ DDL, 95%+ DML, 50%+ stored procedures — with improvements continuously rolled out

Elasticity

with independent
scaling of storage and
compute

Scale up your compute
during peak demand

Scale down your compute
when demand falls

Ramp up your data storage
as you data needs grow

Self-service backup & restore

with up to seven
retained snapshots

Schedule your backups
*to run when it's most convenient for your
business*

Lightening fast restore from snapshot backup,
there when you need it

High Availability

with our robust MPP
architecture

Kubernetes-based container service
*manages all your compute nodes and
automatically responds on node failure*

**Highly-redundant, highly-performance block
storage,** *for all your data*

Db2 Warehouse on Cloud Flex

Our scalable and elastic Cloud data warehouse

***We'll cover a variety of options to
load data in today.***

*Some apply to both Db and Db2
Warehouse on the IBM Cloud. Others will
be product-dependent.*

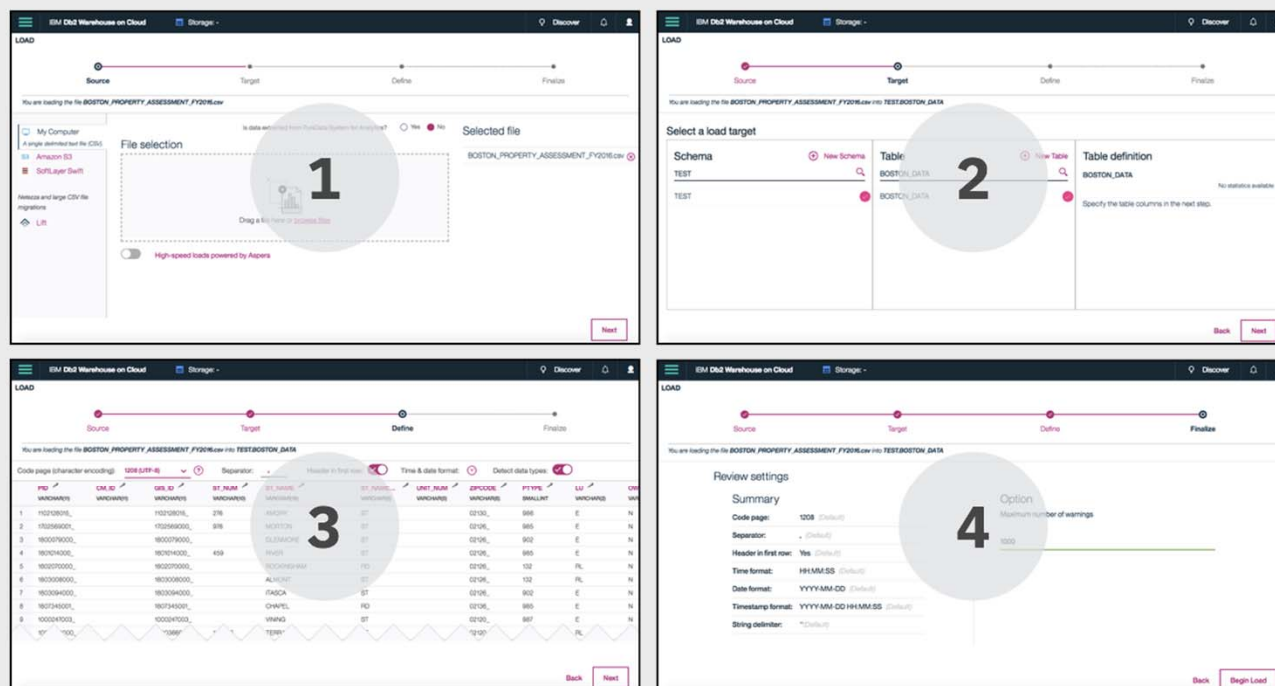
*We'll use insignia to indicate which
method applies to which product.*

Db2 / Db2W

Simple file load through the built-in console

Pros
really easy to go
from a CSV file
data set to running
SQL queries

Cons
one file at a time
file size limits



Choose Load, My Computer and drop a CSV file. Check the column type mapping and hit go.

Pulling data from our IBM Cloud Object Storage with external tables

```
INSERT INTO BOSTON_PROPERTY_ASSESSMENTS
SELECT * FROM EXTERNAL 'BOS_PROP_ASSESSMENTS.csv'
USING(s3('s3-api.us-
geo.objectstorage.service.networklayer.com',
'sJ8uHeupJR4roOFy7NDh',
'xwg72FVrEA47qu09L4tyY3HxksY0sy6yeFZwDzTs',
'property_files_bucket' ));
```

Pros

*programmatically load large data
files into existing tables via SQL
(no additional tooling needed)*

Cons

not point & click, but that's ok

Pulling data from AWS S3

Grab data sets directly from an S3 bucket and ingest into Db2 Warehouse on the IBM Cloud

IBM Db2 Warehouse on Cloud Storage: - Discover

LOAD

Source Target Define Finalize

Select a data source

- My Computer
- S3 Amazon S3**
Delimited text files (CSV, TXT) stored in S3 object storage.
- SoftLayer Swift

Netezza and large CSV file migrations

Lift

Is data extracted from PureData System for Analytics? ☐ Yes ☒ No

File selection

S3 authentication endpoint

US East (N. Virginia) [us-east-1]

S3 access key ID

S3 secret access key

Browse files

Pros

a simple way to liberate your CSV or TXT files by bringing them over to the IBM Cloud

Cons

one file at a time

Over-the-wire migration with the Lift CLI

Lift is...



Blazing-fast

Lift uses IBM Aspera under the covers to move your data to the cloud at blazing fast speeds.



Secure

Nobody wants to end up on the front page of the news. Any data moved over the wire to the IBM Cloud is completely secure via a 256-bit encrypted connection.




Free

Who would charge for data movement, anyway? We want you to try our cloud data services. Cost shouldn't be an issue.



Resilient

Automatically recovers from common problems you'll hit during the migration.



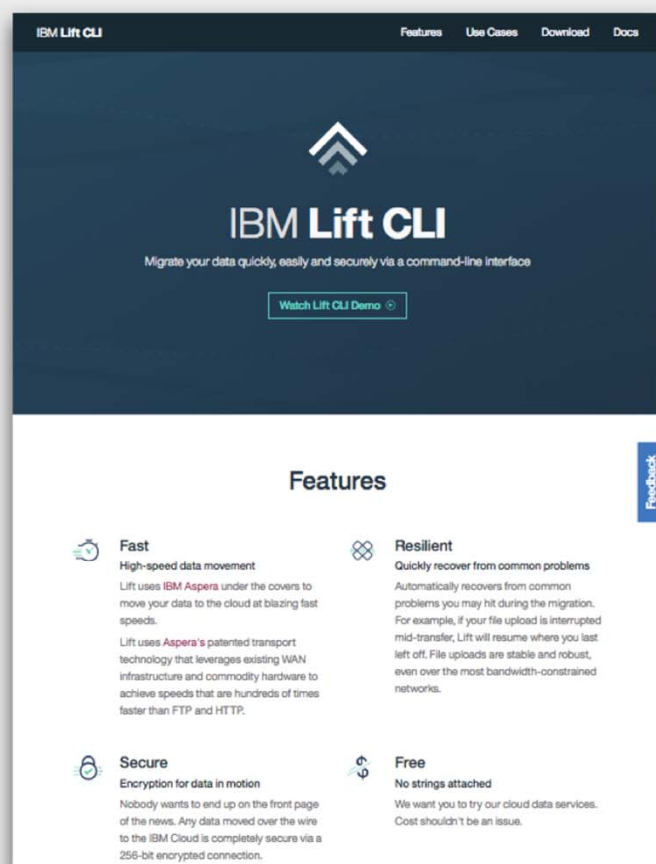
Flexible

Every data migration is split into three steps: extract from source, transport over the wire, and load into target. Run them independently, and on your schedule.



Built for the cloud

You'll install the Lift CLI only once on your on-premises machine. Updates stream automatically.



Migrate data to Db2 or Db2 Warehouse on Cloud in 15 minutes

1. Download the package appropriate to your operating system from ibm.biz/get-lift
2. Unzip the downloaded package
3. Install the package
`% lift-cli-installer/install ~/lift-cli`
4. Extract data from your Netezza database table
`% lift extract -s SCHEMA -t TABLE -f /tmp/TABLE.csv`
5. Transport the file to the Db2 Warehouse on Cloud landing zone
`% lift put -f TABLE.csv`
6. Load the file into the Db2 Warehouse on Cloud engine
`% lift load -f TABLE.csv -s SCHEMA -t TABLE --action replace`

Key Lift CLI capabilities

- **Bulk data migration:**
 - **Source:** CSV files, Netezza/PureData System for Analytics, Integrated Analytics System, Db2, Db2 Warehouse, Oracle
 - **Target:** Db2 & Db2 Warehouse on Cloud
- **High-speed data movement:** Aspera connection with the ability to throttle the speed
- **High-speed extract** on Linux, Windows & macOS via external tables
- **Flexible load options** for Db2 Warehouse on Cloud
- **Over-the-wire data security:** 256-bit encryption for your data traveling over the public Internet
- **Continuous status reporting & end of migration report**
- **Rejected rows** and other logs retrieved automatically for each table
- **Selective extract** with row & column filtering
- **Connection check** route to help get your install ready for data movement
- **Built for the cloud:** new capabilities are automatically rolled out to every instance
- **HIPAA-ready**
- **CLI documentation** with samples

Lift data to Cloud from the IBM Integrated Analytics System

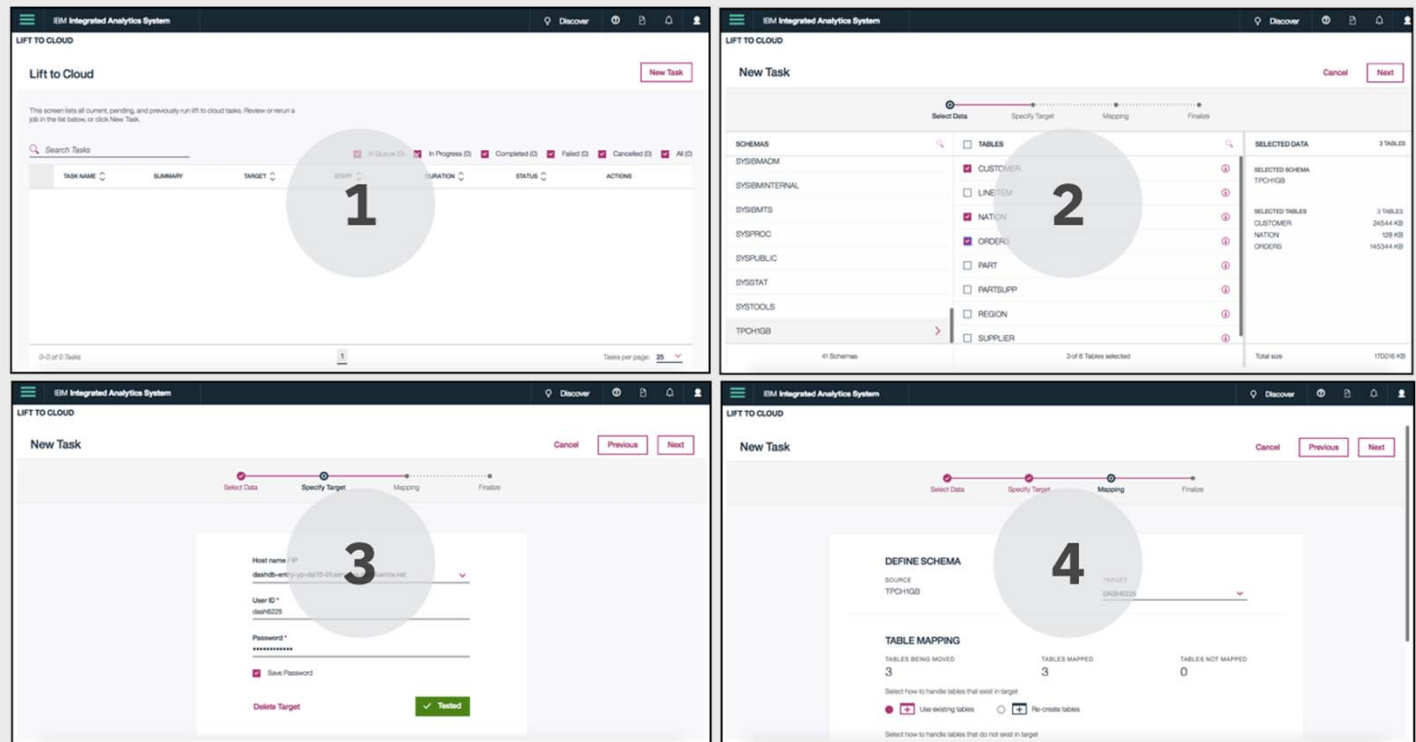
1 Start a new task

2 Pick the schema and tables

3 Provide target connection info

4 Verify table mapping

5 Go



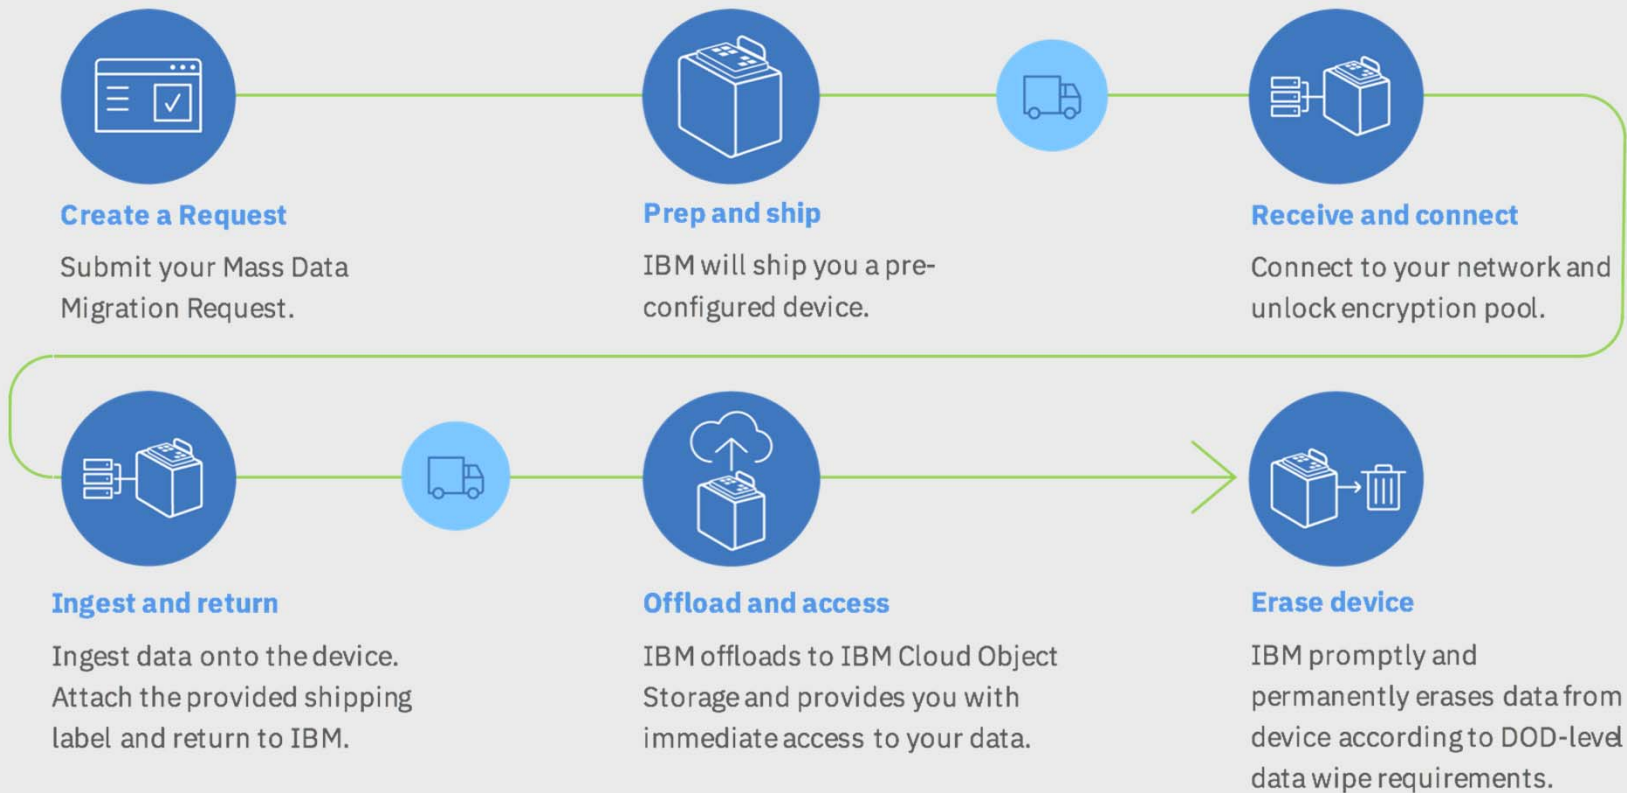
Move large volumes of data to the IBM Cloud with Mass Data Migration

Think 2018 / DOC ID / Month XX, 2018 / © 2018 IBM Corporation


Db2/Db2W



How Mass Data Migration to IBM Cloud works...



How do I decide between Lift and Mass Data Migration service?



Guidance assumes data is uncompressed.

Thank you

Miran Badzak
Senior Product Manager

—
mbadzak@us.ibm.com
ibm.com

